

Overview:

The goal of this lab is to design and integrate a comprehensive set of functions and a pipe structure within the operating system to support the capability of multiprocessing.

Major Parts:

- Synchronization: Add locks for lab1 syscall functions to ensure concurrent access and modification of shared data structures. This will involve the use of semaphores, mutexes or other synchronization mechanisms.
- SysCalls: Implement a set of system call functions, including `fork()`, `wait()`, `exit()` and `exec()` to enable multiprocessing. These functions will be responsible for creating, terminating and managing processes and their execution.
- Pipe: Implement a pipe structure to allow communication between parent and child process. This will involve implementing the `pipe()` function, which creates a new pipe and returns the file descriptors for the read and write ends of the pipe, and the `read()` and `write()` functions for communicating through the pipe.

How different parts of the design interact with each other

`fork()`, `wait()`, and `exit()` will all interact with each other. `fork()` will create new processes and `exit()` will end these processes. `wait()` will be called after `fork()` to wait for child process to `exit()`

The global file struct and `sys_write/sys_read` will interact with the pipe buffer. The pipe buffer will be treated similarly to a file. Where the user will get a file descriptor to read/write to a file buffer.

Synchronization Issues with `fork()`, `wait()`, `exit()`

The goal of `fork()` is for the parent process to make a clone to be able to have a copy of the process. Although the instructions will be the same, the flow of execution can be different depending on the PID, where `fork()` returns the PID of the child process and the child process gets a PID of 0

The goal of `wait()` is for the parent process to halt instruction until any child process finishes. The parent process will clean up child's PCB, page tables, and kernel stack. In the case `wait()` is not called, we will consider how these cases will be handled

The goal of `exit()` is to halt the process and clean up as much resources as possible. The process will let itself be known that it is finished.

In-depth Analysis and Implementation:

We'll to implement `wait()`, `fork()`, `exit()`, `exec()` and pipes:

`wait()`

- Halts the execution of code until any child process is finished executing
- Must acquire spin lock to the ptable when checking child process's state
- Parent cleans up child process on `wait()`
- Need to free child's PCB, page tables, kernel stack
- `wait()` will do a sleep lock to wait for child process
 - Ex: `while (child->state != ZOMBIE) sleep();`
- Returns pid of a finished child process. Returns -1 if there are no child processes (whether all finished or none were called)

`fork()`

- Copy struct context and virtual memory
- Must acquire spin lock to the ptable when creating a new process
- Use `vspaceinstall()` and `vspaceinstallkern()`
- Will have the exact fd pointing to the same `file_struct`. We will have to increase the # of references
- A new entry in the process table must be created via `allocproc`
- User memory must be duplicated via `vspacecopy`
- The trapframe must be duplicated in the new process
- All the opened files must be duplicated in the new process (not as simple as a memory copy)
- Set the state of the new process to be `RUNNABLE`
- Return 0 in the child, while returning the child's pid in the parent

`exit()`

- `exit()` will call `wake()` in order to let parent process know we are finished if they called `wait`
- In the case that `wake()` is called and the parent process never calls `wait()`,
- Set state to ZOMBIE and free up all memory, leave rest to the next process to clean up

`exec()`

- Overwrite current process and start executing the new process
- Create a new virtual address space and register state
- Call `vsploadcode` to read and load programme
- Call `vspaceinitstack` to initialize the user stack
- Call `vspacewriteova` to copy the argument of executing process
- Do not return on success

Pipes

```
struct pipe {
    char *buffer;

    int read_pos;

    int write_pos;

    int max_size;

    struct spinlock lock;
};
```

- Enable intercommunication between parent and child processes
- Runtime allocation using `kalloc()`
- Allocate 4KB as a bounded buffer for the pipe struct
- Free the pipe when the reference counts for both the read end and the write end drop to 0
- Add spinlock to the pipe struct to ensure synchronization
- Add condition variables to the pipe struct to indicate whether the pipe is open for write or read

- Implement `piperead` in `file.c` to read from the pipe buffer as opposed to `concurrent_readi()`
- Implement `pipewrite` in `file.c` to write to the pipe buffer as opposed to `concurrent_writei()`
- Modify `sys_write`, `sys_read`, `sys_dup`, `sys_close` to support pipes

We'll reimplement syscalls from lab1 to support concurrency and pipes:

- Add spinlock to the global file array. Implement the logic in the syscall functions in `sysfile.c` to acquire the lock when they want to access the file array, otherwise, they will get blocked.
- Implement the logic to determine if the passed in file descriptor is a pipe struct. `sys_write` will need to write to the pipe when there is enough space in the buffer, if not the syscall should be blocked until the buffer is freed. That means we require some mechanism, for example, a conditional variable. If the buffer is full, then the process can call `sys_wait` on the conditional variable to block `sys_write`. Or we can use spinlock such that when `sys_write` fail to acquire spinlock, it will be blocked. Same logic applies to `sys_read`, when there is nothing left to read, we will block the `sys_read` until there is data in buffer. For `sys_dup` and `sys_close`, we need to decrement and increment the reference count for our pipe struct.

Risk Analysis:

Unanswered questions

- In the case the parent process does not call `wait()`, do we it up to the next process to clean up?
- What does it mean to clean up a child process? What files do I have to know about?
- How is a pipe implemented? Will we have a `file_struct` in the global file array to write to a pipe or do we need something entirely separate.
- Do we need to add more to our `file_info` struct to be able to hold pipes?

- When to free the lock for pipe buffer when the pipe is filled? How do we free up the space?
- How we distinguish if we are writing or reading to/from a file struct or a pipe?

Staging of work

We'll fix the synchronization of syscalls from lab1. After that we implement `fork()`, `exit()`, `wait()`, `exec()`. Then we will implement the pipe struct. Finally, we implement `sys_pipe()` and modify the syscalls from lab1 to support pipes.

Time estimation

- Lab1 reimplementatation: 4-7 hrs
- `sys_fork()`: 3-5 hrs
- `sys_wait()`: 3-5 hrs
- `sys_exit()`: 3-5 hrs
- `sys_pipe()`: 3-5 hrs
- `sys_exec()`: 3-5 hrs